

cs109a_hw4

October 10, 2017

1 CS 109A/STAT 121A/AC 209A/CSCI E-109A: Homework 4

2 Regularization, High Dimensionality, PCA

Harvard University Fall 2017 Instructors: Pavlos Protopapas, Kevin Rader, Rahul Dave, Margo Levine

2.0.1 INSTRUCTIONS

- To submit your assignment follow the instructions given in canvas.
 - Restart the kernel and run the whole notebook again before you submit.
 - Do not include your name(s) in the notebook even if you are submitting as a group.
 - If you submit individually and you have worked with someone, please include the name of your [one] partner below.
-

Your partner's name (if you submit separately):

Enrollment Status (109A, 121A, 209A, or E109A):

Import libraries:

```
In [12]: import numpy as np
import pandas as pd
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import r2_score
from sklearn.decomposition import PCA
import statsmodels.api as sm
from statsmodels.api import OLS
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.linear_model import RidgeCV
from sklearn.linear_model import LassoCV
%matplotlib inline
sns.set_context("poster")
```

3 Continuing Bike Sharing Usage Data

In this homework, we will focus on multiple linear regression, regularization, dealing with high dimensionality, and PCA. We will continue to build regression models for the Capital Bikeshare program in Washington D.C. See Homework 3 for more information about the data.

*Note: please make sure you use all the processed data from HW 3 Part (a)...you make want to save the data set on your computer and reread the csv/json file here.

```
In [13]: # Create Dataframes
trainDf = pd.read_csv("Bikeshare_train.csv")
testDf = pd.read_csv("Bikeshare_test.csv")

# Create dummy variables
trainBinaryDf = pd.get_dummies(trainDf, columns=["season", "month", "day_of_week", "weather"])
testBinaryDf = pd.get_dummies(testDf, columns=["season", "month", "day_of_week", "weather"])

# Normalize predictors
for column in ["temp", "atemp", "humidity", "windspeed"]:
    colMean = np.mean(trainBinaryDf[column])
    colSD = np.std(trainBinaryDf[column])
    trainBinaryDf[column + "_norm"] = [(x - colMean)/float(colSD) for x in trainBinaryDf[column]]
    testBinaryDf[column + "_norm"] = [(x - colMean)/float(colSD) for x in testBinaryDf[column]]
trainBinaryDf.head()

# Drop k-th dummies
predictors = ["holiday", "workingday", "temp_norm", "atemp_norm", "humidity_norm", "windspeed_norm",
              "weather_1.0", "weather_2.0", "day_of_week_0.0", "day_of_week_1.0", "day_of_week_2.0",
              "day_of_week_3.0", "day_of_week_4.0", "day_of_week_5.0", "season_1.0", "season_2.0",
              "season_3.0", "month_2.0", "month_3.0", "month_4.0", "month_5.0", "month_6.0",
              "month_7.0", "month_8.0", "month_9.0", "month_10.0", "month_11.0"]

# Add a constant term to our predictors
trainBinaryDf = sm.add_constant(trainBinaryDf)
testBinaryDf = sm.add_constant(testBinaryDf)

trainPredict = trainBinaryDf[predictors]
testPredict = testBinaryDf[predictors]
```

In the above code, we simply prepared our data in the same way that we did in HW3.

3.1 Part (f): Regularization/Penalization Methods

As an alternative to selecting a subset of predictors and fitting a regression model on the subset, one can fit a linear regression model on all predictors, but shrink or regularize the coefficient estimates to make sure that the model does not "overfit" the training set.

Use the following regularization techniques to fit linear models to the training set: - Ridge regression - Lasso regression

You may choose the shrinkage parameter λ from the set $\{10^{-5}, 10^{-4}, \dots, 10^4, 10^5\}$ using cross-validation. In each case,

- How do the estimated coefficients compare to or differ from the coefficients estimated by a plain linear regression (without shrinkage penalty) in Part (b) from HW 3? Is there a difference between coefficients estimated by the two shrinkage methods? If so, give an explanation for the difference.
- List the predictors that are assigned a coefficient value close to 0 (say $< 1e-10$) by the two methods. How closely do these predictors match the redundant predictors (if any) identified in Part (c) from HW 3?
- Is there a difference in the way Ridge and Lasso regression assign coefficients to the predictors `temp` and `atemp`? If so, explain the reason for the difference.

We next analyze the performance of the two shrinkage methods for different training sample sizes: - Generate random samples of sizes 100, 150, ..., 400 from the training set. You may use the following code to draw a random sample of a specified size from the training set:

```
In [14]: #----- sample
# A function to select a random sample of size k from the training set
# Input:
#     x (n x d array of predictors in training data)
#     y (n x 1 array of response variable vals in training data)
#     k (size of sample)
# Return:
#     chosen sample of predictors and responses

def sample(x, y, k):
    n = x.shape[0] # No. of training points

    # Choose random indices of size 'k'
    subset_ind = np.random.choice(np.arange(n), k)

    # Get predictors and responses with the indices
    x_subset = x[subset_ind, :]
    y_subset = y[subset_ind]

    return (x_subset, y_subset)
```

In []:

- Fit linear, Ridge and Lasso regression models to each of the generated sample. In each case, compute the R^2 score for the model on the training sample on which it was fitted, and on the test set.
- Repeat the above experiment for 10 random trials/splits, and compute the average train and test R^2 across the trials for each training sample size. Also, compute the standard deviation (SD) in each case.
- Make a plot of the mean training R^2 scores for the linear, Ridge and Lasso regression methods as a function of the training sample size. Also, show a confidence interval for the mean scores extending from **mean - SD** to **mean + SD**. Make a similar plot for the test R^2 scores.

How do the training and test R^2 scores compare for the three methods? Give an explanation for your observations. How do the confidence intervals for the estimated R^2 change with training

sample size? Based on the plots, which of the three methods would you recommend when one needs to fit a regression model using a small training sample?

Hint: You may use sklearn's RidgeCV and LassoCV classes to implement Ridge and Lasso regression. These classes automatically perform cross-validation to tune the parameter λ from a given range of values. You may use the plt.errorbar function to plot confidence bars for the average R^2 scores.

```
In [15]: # Lambdas
         lambdas = [10**x for x in range(-5, 6)]
         xTrain = trainPredict
         yTrain = trainBinaryDf["count"]

         ridgeCoefs = []
         lassoCoefs = []

         # Regression on entire training data set
         ridge = RidgeCV(alphas=lambdas, cv=10, fit_intercept=True)
         ridge.fit(xTrain, yTrain)

         ridgeCoefs = ridge.coef_

         lasso = LassoCV(alphas=lambdas, cv=10, max_iter=10000, fit_intercept=True)
         lasso.fit(xTrain, yTrain)

         lassoCoefs = lasso.coef_

In [16]: HW3Model = sm.OLS(yTrain, xTrain).fit()
         HW3Coefs = HW3Model.params

In [17]: r2TrainVals = pd.DataFrame(data={"Original": [.576], "Lasso": [lasso.score(xTrain, yTrain)],
                                         "Ridge": [ridge.score(xTrain, yTrain)]})
         r2TrainVals

Out[17]:
```

	Lasso	Original	Ridge
0	0.563605	0.576	0.559563

```
In [18]: xTest = testPredict
         yTest = testBinaryDf["count"]

In [19]: r2TestVals = pd.DataFrame(data={"Original": [.249], "Lasso": [lasso.score(xTest, yTest)],
                                         "Ridge": [ridge.score(xTest, yTest)]})
         r2TestVals

Out[19]:
```

	Lasso	Original	Ridge
0	0.263547	0.249	0.256219

```
In [20]: ridgeCoefs[-1] = ridge.intercept_
         lassoCoefs[-1] = lasso.intercept_
         coefsTable = pd.DataFrame(data={"HW3Model": HW3Coefs, "RidgeModel": ridgeCoefs, "LassoModel": lassoCoefs})
         coefsTable.head(50)
```

Out [20]:

	HW3Model	LassoModel	RidgeModel
holiday	-616.602710	-171.307907	-235.657032
workingday	-24.093294	0.000000	28.943832
temp_norm	924.334403	615.501808	596.419466
atemp_norm	311.961760	473.093575	531.323759
humidity_norm	-547.663783	-556.194283	-558.385010
windspeed_norm	-254.736916	-246.996085	-264.834036
weather_1.0	1581.978284	759.077243	480.235213
weather_2.0	1565.411700	731.541784	452.456290
day_of_week_0.0	-465.145010	-293.747624	-268.575304
day_of_week_1.0	-256.650051	-204.727796	-203.553159
day_of_week_2.0	-328.184507	-141.402923	-169.398103
day_of_week_3.0	37.612773	2.027548	67.202563
day_of_week_4.0	-71.642544	-0.000000	15.123090
day_of_week_5.0	-21.831675	19.958391	83.912409
season_1.0	-1226.186543	-980.105861	-831.421968
season_2.0	-327.357503	-0.000000	-71.167428
season_3.0	-193.304968	-225.728921	-236.541870
month_2.0	88.940093	0.000000	-43.283473
month_3.0	239.180898	21.848502	96.419673
month_4.0	333.349086	105.307226	243.948855
month_5.0	-65.812500	-0.000000	57.943527
month_6.0	-792.262899	-409.762166	-357.529453
month_7.0	-1279.987006	-633.987086	-578.242146
month_8.0	-776.475490	-79.491851	-151.051370
month_9.0	405.144566	743.960587	640.829392
month_10.0	486.250904	673.115269	607.588625
month_11.0	112.681645	114.448501	192.866358
month_12.0	-118.835819	-0.000000	-60.494888
const	3791.129801	4189.010483	4401.735647

In [21]: `def randSampleR2(k):`

"""k is the sample size

outputs R2 values for linear, lasso, ridge regression models"""

`xSub, ySub = sample(xTrain.as_matrix(), yTrain.as_matrix(), k)`

`ridge = RidgeCV(alphas=lambdas, cv=10, fit_intercept=True)`

`ridge.fit(xSub, ySub)`

`lasso = LassoCV(alphas=lambdas, cv=10, fit_intercept=True, max_iter=10000)`

`lasso.fit(xSub, ySub)`

`linear = sm.OLS(ySub, xSub).fit()`

`yHat = linear.predict(xTest)`

```

        return linear.rsquared, lasso.score(xSub, ySub), ridge.score(xSub, ySub), \
               r2_score(yTest, yHat), lasso.score(xTest, yTest), ridge.score(xTest, yTest)

In [22]: sampleSizes = np.arange(100, 450, 50)

rSquaredBootStrap = {"sampleSize": sampleSizes, "linearTrain": [], "lassoTrain": [],
                    "ridgeTrain": [], "linearTest": [], "lassoTest": [], "ridgeTest": []}

for k in sampleSizes:
    linearTestR2 = []
    lassoTestR2 = []
    ridgeTestR2 = []
    linearTrainR2 = []
    lassoTrainR2 = []
    ridgeTrainR2 = []
    for _ in range(10):
        linTrain, lasTrain, ridTrain, linTest, lasTest, ridTest = randSampleR2(k)
        linearTrainR2.append(linTrain)
        lassoTrainR2.append(lasTrain)
        ridgeTrainR2.append(ridTrain)
        linearTestR2.append(linTest)
        lassoTestR2.append(lasTest)
        ridgeTestR2.append(ridTest)

    rSquaredBootStrap["linearTrain"].append(linearTrainR2)
    rSquaredBootStrap["lassoTrain"].append(lassoTrainR2)
    rSquaredBootStrap["ridgeTrain"].append(ridgeTrainR2)
    rSquaredBootStrap["linearTest"].append(linearTestR2)
    rSquaredBootStrap["lassoTest"].append(lassoTestR2)
    rSquaredBootStrap["ridgeTest"].append(ridgeTestR2)

In [23]: fig, ax = plt.subplots()

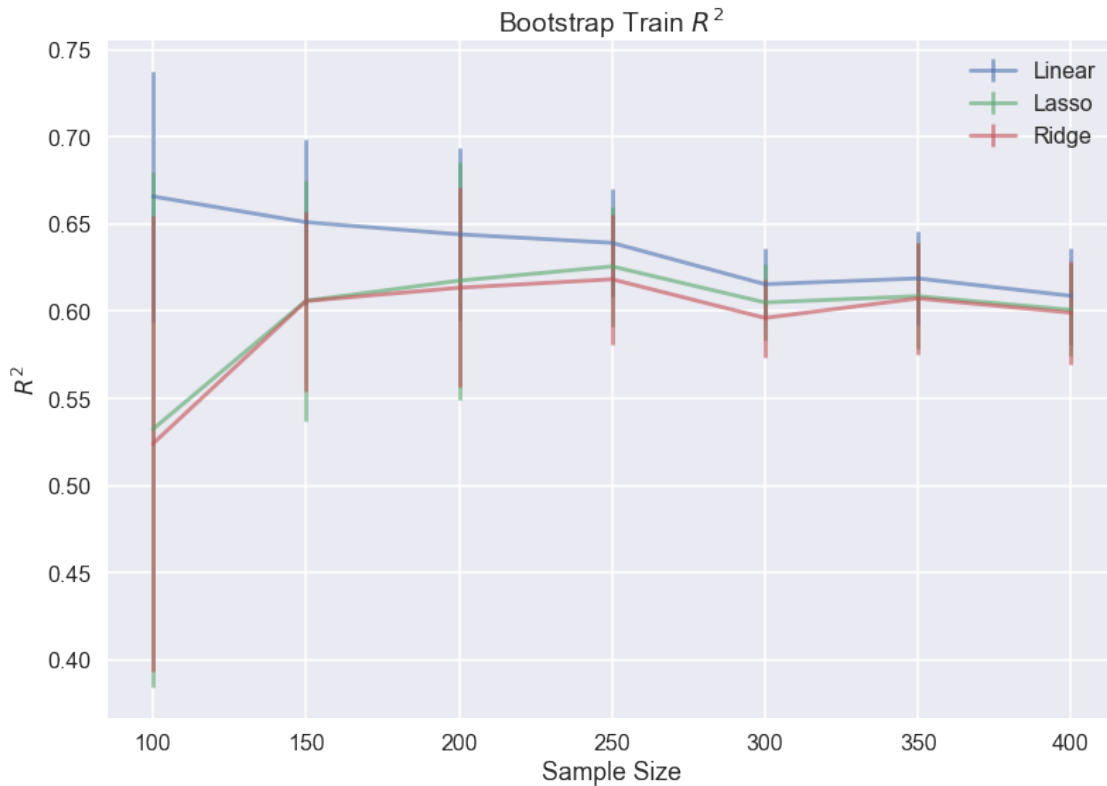
linTrainMeans = [np.mean(x) for x in rSquaredBootStrap["linearTrain"]]
linTrainSDs = [np.std(x) for x in rSquaredBootStrap["linearTrain"]]
lasTrainMeans = [np.mean(x) for x in rSquaredBootStrap["lassoTrain"]]
lasTrainSDs = [np.std(x) for x in rSquaredBootStrap["lassoTrain"]]
ridTrainMeans = [np.mean(x) for x in rSquaredBootStrap["ridgeTrain"]]
ridTrainSDs = [np.std(x) for x in rSquaredBootStrap["ridgeTrain"]]

ax.set_title("Bootstrap Train  $R^2$ ")
ax.set_xlabel("Sample Size")
ax.set_ylabel(" $R^2$ ")

ax.errorbar(rSquaredBootStrap["sampleSize"], linTrainMeans, linTrainSDs, label="Linear")
ax.errorbar(rSquaredBootStrap["sampleSize"], lasTrainMeans, lasTrainSDs, label="Lasso")
ax.errorbar(rSquaredBootStrap["sampleSize"], ridTrainMeans, ridTrainSDs, label="Ridge")
ax.legend()

Out[23]: <matplotlib.legend.Legend at 0x122fb93c8>

```



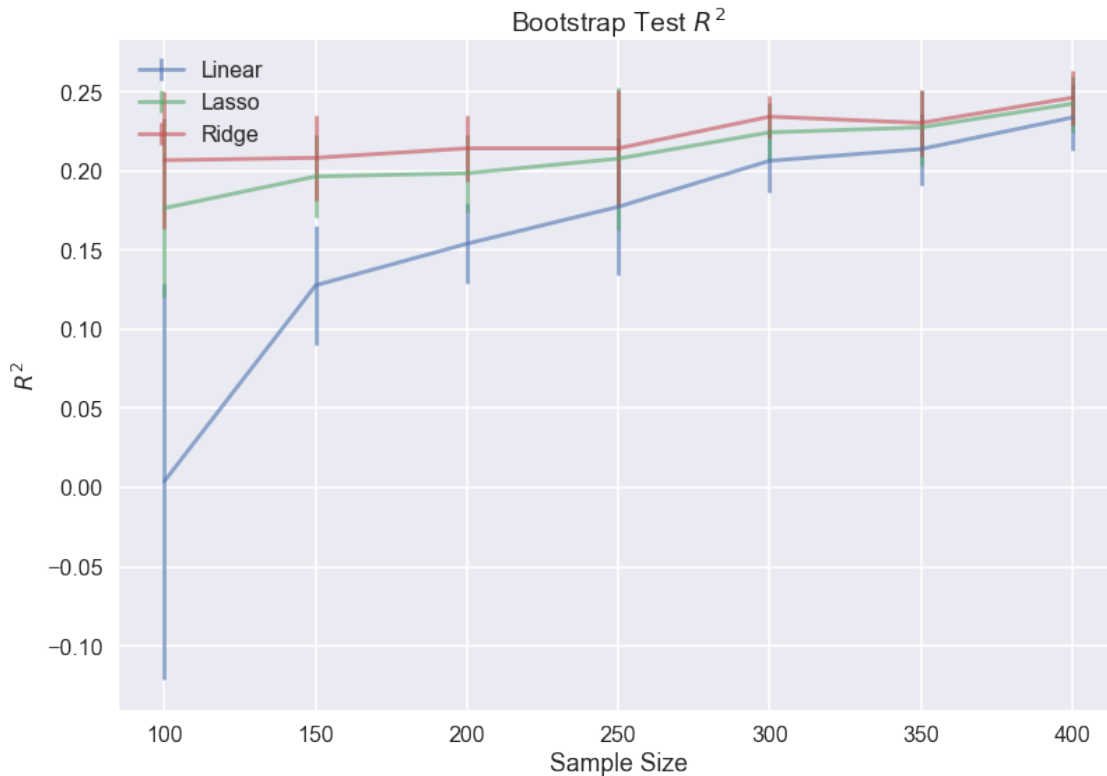
```
In [24]: fig, ax = plt.subplots()
```

```
linTestMeans = [np.mean(x) for x in rSquaredBootStrap["linearTest"]]
linTestSDs = [np.std(x) for x in rSquaredBootStrap["linearTest"]]
lasTestMeans = [np.mean(x) for x in rSquaredBootStrap["lassoTest"]]
lasTestSDs = [np.std(x) for x in rSquaredBootStrap["lassoTest"]]
ridTestMeans = [np.mean(x) for x in rSquaredBootStrap["ridgeTest"]]
ridTestSDs = [np.std(x) for x in rSquaredBootStrap["ridgeTest"]]
```

```
ax.set_title("Bootstrap Test  $R^2$ ")
ax.set_xlabel("Sample Size")
ax.set_ylabel(" $R^2$ ")
```

```
ax.errorbar(rSquaredBootStrap["sampleSize"], linTestMeans, linTestSDs, label="Linear",
ax.errorbar(rSquaredBootStrap["sampleSize"], lasTestMeans, lasTestSDs, label="Lasso", a
ax.errorbar(rSquaredBootStrap["sampleSize"], ridTestMeans, ridTestSDs, label="Ridge", a
ax.legend()
```

```
Out[24]: <matplotlib.legend.Legend at 0x1232bae10>
```



3.2 Response to above questions

- How do the estimated coefficients compare to or differ from the coefficients estimated by a plain linear regression (without shrinkage penalty) in Part (b) from HW 3? Is there a difference between coefficients estimated by the two shrinkage methods? If so, give an explanation for the difference.

As we can see from the data table above, the magnitude of our Lasso and Ridge coefficients are significantly and consistently smaller than those of our HW3 model. This is exactly the behavior we would expect, as Lasso and Ridge methods try to decrease the magnitude of our coefficients by adding a cost term for the L1 or L2 of our β .

- List the predictors that are assigned a coefficient value close to 0 (say $< 1e-10$) by the two methods. How closely do these predictors match the redundant predictors (if any) identified in Part (c) from HW 3?

Lasso Model: *working_day*, *day_of_week_4.0*, *season_2.0*, *month_2.0*, *month_5.0*, *month_12.0*. We note that there are no coefficient values close to 0 in either Ridge or our HW3 models. This makes sense, as shown in class, that it is more likely for the L1 norm cost function to result in coefficients of value 0 than for L2 norm.

- Is there a difference in the way Ridge and Lasso regression assign coefficients to the predictors *temp* and *atemp*? If so, explain the reason for the difference.

Both Ridge and Lasso regression preserves the initial relationship observed in HW3, that *temp* has a larger coefficient value than *atemp*. However, what is interesting to note is that the difference between these coefficients seem to converge to 0 as our cost term increases its L^i dimensional norm. For example, for Lasso in L^1 , the difference is much larger than for Ridge, which uses L^2 .

- How do the training and test R^2 scores compare for the three methods? Give an explanation for your observations. How do the confidence intervals for the estimated R^2 change with training sample size? Based on the plots, which of the three methods would you recommend when one needs to fit a regression model using a small training sample?

We notice that the training R^2 values for both Lasso and Ridge are lower than our HW3 model, which is to be expected as we are introducing a regularization parameter. However, our Lasso and Ridge models perform better on the test data than our HW3 model, implying that we have successfully mitigated some of the overfitting from our HW3 model.

The confidence intervals for the estimated R^2 tend to decrease as our sample size grows. While there are some fluctuations in the above graphs, we note that this overall trend of decreasing confidence intervals seems to dominate over all three models.

If we had to choose one of the three models given a small training sample, we would select the Ridge regression model. It has the smallest standard deviation in its R^2 scores, and consistently outperforms the other two models over many small sample sizes.

3.3 Part (g): Polynomial & Interaction Terms

Moving beyond linear models, we will now try to improve the performance of the regression model in Part (b) from HW 3 by including higher-order polynomial and interaction terms.

- For each continuous predictor X_j , include additional polynomial terms X_j^2 , X_j^3 , and X_j^4 , and fit a multiple regression model to the expanded training set. How does the R^2 of this model on the test set compare with that of the linear model fitted in Part (b) from HW 3? Using a t-test, find out which of estimated coefficients for the polynomial terms are statistically significant at a significance level of 5%.
- Fit a multiple linear regression model with additional interaction terms $\mathbb{I}_{month=12} \times temp$ and $\mathbb{I}_{workingday=1} \times \mathbb{I}_{weathersit=1}$ and report the test R^2 for the fitted model. How does this compare with the R^2 obtained using linear model in Part (b) from HW 3? Are the estimated coefficients for the interaction terms statistically significant at a significance level of 5%?

```
In [25]: # Generate polynomial features
multTrainPredict = trainPredict.copy()
multTestPredict = testPredict.copy()

for i in range(2,5):
    for var in ["temp_norm", "atemp_norm", "humidity_norm", "windspeed_norm"]:
        multTrainPredict["{}_{}".format(var, i)] = multTrainPredict[var]**i
        multTestPredict["{}_{}".format(var, i)] = multTestPredict[var]**i

In [26]: # Fit OLS model
multModel = sm.OLS(yTrain, multTrainPredict).fit()
multModel.summary()
```

```
Out[26]: <class 'statsmodels.iolib.summary.Summary'>
        """
```

OLS Regression Results

Dep. Variable:	count	R-squared:	0.670			
Model:	OLS	Adj. R-squared:	0.625			
Method:	Least Squares	F-statistic:	15.13			
Date:	Tue, 10 Oct 2017	Prob (F-statistic):	7.98e-50			
Time:	13:29:19	Log-Likelihood:	-2790.9			
No. Observations:	331	AIC:	5662.			
Df Residuals:	291	BIC:	5814.			
Df Model:	39					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

holiday	-526.2557	372.773	-1.412	0.159	-1259.929	207.418
workingday	14.7858	159.059	0.093	0.926	-298.265	327.837
temp_norm	770.3204	758.967	1.015	0.311	-723.441	2264.082
atemp_norm	895.9192	712.094	1.258	0.209	-505.588	2297.426
humidity_norm	-667.9033	157.118	-4.251	0.000	-977.136	-358.671
windspeed_norm	-445.8335	148.704	-2.998	0.003	-738.505	-153.162
weather_1.0	1043.9997	546.051	1.912	0.057	-30.710	2118.709
weather_2.0	1103.0116	499.023	2.210	0.028	120.859	2085.164
day_of_week_0.0	-471.0834	246.557	-1.911	0.057	-956.345	14.178
day_of_week_1.0	-227.9217	157.372	-1.448	0.149	-537.653	81.809
day_of_week_2.0	-268.0232	186.488	-1.437	0.152	-635.060	99.014
day_of_week_3.0	13.1360	196.966	0.067	0.947	-374.522	400.794
day_of_week_4.0	-104.0028	189.352	-0.549	0.583	-476.675	268.669
day_of_week_5.0	75.3419	184.193	0.409	0.683	-287.178	437.861
season_1.0	-1523.2288	467.580	-3.258	0.001	-2443.496	-602.961
season_2.0	-756.7981	536.808	-1.410	0.160	-1813.316	299.720
season_3.0	55.5255	419.974	0.132	0.895	-771.046	882.097
month_2.0	-325.0686	409.611	-0.794	0.428	-1131.245	481.108
month_3.0	-304.8491	446.028	-0.683	0.495	-1182.700	573.002
month_4.0	-418.0245	639.524	-0.654	0.514	-1676.703	840.654
month_5.0	-1037.2042	677.186	-1.532	0.127	-2370.008	295.599
month_6.0	-1456.1857	697.520	-2.088	0.038	-2829.010	-83.362
month_7.0	-1416.9882	749.751	-1.890	0.060	-2892.610	58.634
month_8.0	-1715.9389	743.240	-2.309	0.022	-3178.747	-253.131
month_9.0	-1073.4008	660.859	-1.624	0.105	-2374.069	227.268
month_10.0	-925.8710	617.522	-1.499	0.135	-2141.247	289.505
month_11.0	-825.5328	591.138	-1.397	0.164	-1988.981	337.916
month_12.0	-555.6676	479.543	-1.159	0.248	-1499.481	388.146
const	5985.5838	836.335	7.157	0.000	4339.552	7631.616
temp_norm_2	-1805.5466	814.442	-2.217	0.027	-3408.491	-202.602
atemp_norm_2	1171.9491	786.481	1.490	0.137	-375.962	2719.861
humidity_norm_2	-53.5087	154.914	-0.345	0.730	-358.402	251.384

windspeed_norm_2	-34.0621	126.569	-0.269	0.788	-283.168	215.044
temp_norm_3	8.5688	274.482	0.031	0.975	-531.653	548.790
atemp_norm_3	-302.5595	244.983	-1.235	0.218	-784.722	179.603
humidity_norm_3	-15.9849	44.689	-0.358	0.721	-103.940	71.970
windspeed_norm_3	44.6309	65.163	0.685	0.494	-83.619	172.881
temp_norm_4	-44.9184	170.385	-0.264	0.792	-380.261	290.425
atemp_norm_4	-20.6433	146.714	-0.141	0.888	-309.398	268.112
humidity_norm_4	-24.6869	31.291	-0.789	0.431	-86.272	36.898
windspeed_norm_4	-20.0552	30.144	-0.665	0.506	-79.382	39.272

```
=====
Omnibus:                29.995    Durbin-Watson:                1.959
Prob(Omnibus):           0.000    Jarque-Bera (JB):         10.202
Skew:                   -0.094    Prob(JB):                 0.00609
Kurtosis:                2.161    Cond. No.                  3.54e+16
=====
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The smallest eigenvalue is 3.1e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.
"""
```

```
In [27]: yHatTest = multModel.predict(multTestPredict)
        r2_score(yTest, yHatTest)
```

```
Out[27]: 0.27723843508616397
```

```
In [28]: # Generate additional interaction features
        intTrainPredict = multTrainPredict.copy()
        intTestPredict = multTestPredict.copy()

        # Add month=12 x temp interaction terms
        intTrainPredict["month_12xtemp_norm"] = intTrainPredict["month_12.0"] * intTrainPredict["temp_norm"]
        intTestPredict["month_12xtemp_norm"] = intTestPredict["month_12.0"] * intTestPredict["temp_norm"]

        # Add workingday x weathersit=1 interaction terms
        intTrainPredict["workingdayxweather_1"] = intTrainPredict["workingday"] * intTrainPredict["weathersit"]
        intTestPredict["workingdayxweather_1"] = intTestPredict["workingday"] * intTestPredict["weathersit"]
```

```
In [29]: # Fit Interaction OLS model
        intModel = sm.OLS(yTrain, intTrainPredict).fit()
        intModel.summary()
```

```
Out[29]: <class 'statsmodels.iolib.summary.Summary'>
        """
```

```

                                OLS Regression Results
=====
Dep. Variable:                count    R-squared:                0.670
Model:                        OLS    Adj. R-squared:            0.623

```

Method:	Least Squares	F-statistic:	14.31
Date:	Tue, 10 Oct 2017	Prob (F-statistic):	1.07e-48
Time:	13:29:19	Log-Likelihood:	-2790.7
No. Observations:	331	AIC:	5665.
Df Residuals:	289	BIC:	5825.
Df Model:	41		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975
-----	-----	-----	-----	-----	-----	-----
holiday	-504.8692	376.774	-1.340	0.181	-1246.437	236.69
workingday	-78.3221	242.127	-0.323	0.747	-554.878	398.23
temp_norm	793.8768	763.389	1.040	0.299	-708.631	2296.38
atemp_norm	880.3824	715.945	1.230	0.220	-528.745	2289.51
humidity_norm	-675.3295	158.264	-4.267	0.000	-986.826	-363.83
windspeed_norm	-446.4038	149.948	-2.977	0.003	-741.532	-151.27
weather_1.0	923.0406	594.656	1.552	0.122	-247.366	2093.44
weather_2.0	1104.9141	502.004	2.201	0.029	116.866	2092.96
day_of_week_0.0	-462.4486	248.640	-1.860	0.064	-951.824	26.92
day_of_week_1.0	-239.5854	159.457	-1.503	0.134	-553.429	74.25
day_of_week_2.0	-280.3233	188.566	-1.487	0.138	-651.459	90.81
day_of_week_3.0	-3.0325	200.247	-0.015	0.988	-397.160	391.09
day_of_week_4.0	-119.9933	192.772	-0.622	0.534	-499.408	259.42
day_of_week_5.0	59.7432	187.628	0.318	0.750	-309.547	429.03
season_1.0	-1493.1626	505.353	-2.955	0.003	-2487.801	-498.52
season_2.0	-729.5247	549.124	-1.329	0.185	-1810.315	351.26
season_3.0	71.1170	422.825	0.168	0.867	-761.090	903.32
month_2.0	-325.7395	415.794	-0.783	0.434	-1144.107	492.62
month_3.0	-311.7359	465.982	-0.669	0.504	-1228.885	605.41
month_4.0	-429.5823	676.680	-0.635	0.526	-1761.428	902.26
month_5.0	-1024.8594	711.364	-1.441	0.151	-2424.970	375.25
month_6.0	-1466.6541	734.348	-1.997	0.047	-2912.002	-21.30
month_7.0	-1427.6739	798.423	-1.788	0.075	-2999.136	143.78
month_8.0	-1720.6375	791.728	-2.173	0.031	-3278.922	-162.35
month_9.0	-1055.9093	726.079	-1.454	0.147	-2484.983	373.16
month_10.0	-882.2336	692.220	-1.274	0.204	-2244.665	480.19
month_11.0	-773.6133	656.762	-1.178	0.240	-2066.257	519.03
month_12.0	-480.9222	1027.332	-0.468	0.640	-2502.924	1541.08
const	6035.4428	902.323	6.689	0.000	4259.486	7811.40
temp_norm_2	-1796.4144	817.390	-2.198	0.029	-3405.207	-187.62
atemp_norm_2	1167.1361	790.067	1.477	0.141	-387.880	2722.15
humidity_norm_2	-56.3479	155.509	-0.362	0.717	-362.421	249.72
windspeed_norm_2	-37.7429	127.475	-0.296	0.767	-288.641	213.15
temp_norm_3	1.6553	277.479	0.006	0.995	-544.481	547.79
atemp_norm_3	-297.6352	246.540	-1.207	0.228	-782.877	187.60
humidity_norm_3	-14.9630	44.946	-0.333	0.739	-103.425	73.49
windspeed_norm_3	44.3740	65.415	0.678	0.498	-84.376	173.12
temp_norm_4	-44.4857	170.895	-0.260	0.795	-380.843	291.87

atemp_norm_4	-20.9201	147.604	-0.142	0.887	-311.435	269.59
humidity_norm_4	-23.8429	31.490	-0.757	0.450	-85.822	38.13
windspeed_norm_4	-19.2421	30.292	-0.635	0.526	-78.862	40.37
month_12xtemp_norm	44.3435	714.542	0.062	0.951	-1362.024	1450.71
workingdayxweather_1	165.0265	322.219	0.512	0.609	-469.167	799.22

Omnibus:	29.023	Durbin-Watson:	1.960
Prob(Omnibus):	0.000	Jarque-Bera (JB):	10.080
Skew:	-0.099	Prob(JB):	0.00647
Kurtosis:	2.168	Cond. No.	3.92e+16

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified
[2] The smallest eigenvalue is 2.52e-29. This might indicate that there are
strong multicollinearity problems or that the design matrix is singular.
"""
```

```
In [30]: yHatTest = intModel.predict(intTestPredict)
r2_score(yTest, yHatTest)
```

```
Out [30]: 0.28297793302407059
```

3.4 Answers to Above Questions

- For each continuous predictor X_j , include additional polynomial terms X_j^2 , X_j^3 , and X_j^4 , and fit a multiple regression model to the expanded training set. How does the R^2 of this model on the test set compare with that of the linear model fitted in Part (b) from HW 3? Using a t-test, find out which of estimated coefficients for the polynomial terms are statistically significant at a significance level of 5%.

We found that the R^2 of this model was .277 for the test data set, which clearly outperforms the R^2 score of our HW3 model on the test data set of .249. Using the t-tests provided in the summary above, we found that the only statistically significant coefficient was $temp_norm_2$, which equates to X^2 where X is the normalized continuous variable for temperature.

- Fit a multiple linear regression model with additional interaction terms $\mathbb{I}_{month=12} \times temp$ and $\mathbb{I}_{workingday=1} \times \mathbb{I}_{weathersit=1}$ and report the test R^2 for the fitted model. How does this compare with the R^2 obtained using linear model in Part (b) from HW 3? Are the estimated coefficients for the interaction terms statistically significant at a significance level of 5%?

For this additional model, we found that the R^2 score on the test set was .283, even better than our previous model of .277, and much better than the HW3 model of .249. There are no statistically significant interaction coefficients, as their 95% confidence intervals, as shown above, include 0.

3.5 Part (h): PCA to deal with high dimensionality

We would like to fit a model to include all main effects, polynomial terms up to the 4th order, and all interactions between all possible predictors and polynomial terms (not including the interactions between X_j^1 , X_j^2 , X_j^3 , and X_j^4 as they would just create higher order polynomial terms).

- Create an expanded training set including all the desired terms mentioned above. What are the dimensions of this 'design matrix' of all the predictor variables? What are the issues with attempting to fit a regression model using all of these predictors?
- Instead of using the usual approaches for model selection, let's instead use principal components analysis (PCA) to fit the model. First, create the principal component vectors in python (consider: should you normalize first?). Then fit 5 different regression models: (1) using just the first PCA vector, (2) using the first two PCA vectors, (3) using the first three PCA vectors, etc... Briefly summarize how these models compare in the training set.
- Use the test set to decide which of the 5 models above is best to predict out of sample. How does this model compare to the previous models you've fit? What are the interpretations of this model's coefficients?

```
In [31]: xTest = testPredict
        yTest = testBinaryDf["count"]

        xTrain = trainPredict
        yTrain = trainBinaryDf["count"]

        genPolyTerms = PolynomialFeatures(degree=4, interaction_only=False)

        xTrainPoly = []
        xTestPoly = []

        xTrainPoly.append(genPolyTerms.fit_transform(xTrain))
        xTestPoly.append(genPolyTerms.fit_transform(xTest))

        print("There are {} predictors in this new design matrix.".format(xTrainPoly[0].shape[1]))
```

There are 40920 predictors in this new design matrix.

```
In [32]: pcaXTrainVals = []
        pcaXTestVals = []
        for i in range(1,6):
            pca = PCA(n_components=i)
            pca.fit(xTrainPoly[0])
            xPCA = pca.transform(xTrainPoly[0])
            xTestPCA = pca.transform(xTestPoly[0])
            pcaXTrainVals.append(xPCA)
            pcaXTestVals.append(xTestPCA)
```

```
In [33]: trainPCAModels = []
        r2sTrain = []
        r2sTest = []

        for i, data in enumerate(pcaXTrainVals):
            regPCAModel = LinearRegression(fit_intercept=True)
```

```

regPCAModel.fit(data, yTrain)
r2sTrain.append(regPCAModel.score(data, yTrain))
r2sTest.append(regPCAModel.score(pcaXTestVals[i], yTest))
trainPCAModels.append(regPCAModel)

```

```

In [34]: fig, ax = plt.subplots()
ax.plot(range(1,6), r2sTrain, label="Train  $R^2$ ")
ax.plot(range(1,6), r2sTest, label="Test  $R^2$ ")
ax.set_xlabel("Number of Principle Components")
ax.set_ylabel(" $R^2$ ")
ax.set_title("Comparison of Training vs Test  $R^2$  for PCA Analysis")
plt.xticks(range(1,6), )
ax.legend()


```

r2sTrain, r2sTest

```

Out [34]: ([0.12699183896138022,
0.17662017642476535,
0.24654449750664345,
0.28015860595895636,
0.34565321930651116],
[0.07156378309358824,
0.14324347551193373,
0.17721903027072172,
0.14923996544624762,
0.13943327928477645])

```



```
In [35]: pcaGoalModel = trainPCAModels[2]
         print("The coefficients are: {}".format(pcaGoalModel.coef_))
         print("The intercept is: {}".format(pcaGoalModel.intercept_))
```

```
The coefficients are: [-34.95509138 -32.72310621  48.22779015].
The intercept is: 4598.447129909366.
```

3.6 Answers to Above Questions

- Create an expanded training set including all the desired terms mentioned above. What are the dimensions of this 'design matrix' of all the predictor variables? What are the issues with attempting to fit a regression model using all of these predictors?

We can see that there are 40,920 predictor variables in this new design matrix, with 331 rows. One major reason why we would not want to fit a model with this many predictors is because the number of predictors is larger than the number of data points, implying that our model will completely overfit our data. Additionally, we will have collinearity, and model selection (applying stepwise forward algorithms) become intractable.

- Instead of using the usual approaches for model selection, let's instead use principal components analysis (PCA) to fit the model. First, create the principal component vectors in python (consider: should you normalize first?). Then fit 5 different regression models: (1) using just the first PCA vector, (2) using the first two PCA vectors, (3) using the first three PCA vectors, etc... Briefly summarize how these models compare in the training set.

From the above plot, we can see that as the number of principle components included in our model increases, our training R^2 monotonically increases. This makes sense, as we are doing a better job at capturing the variance in our training data with each additional principle component. However, for the test data, we find that we begin to overfit our training data after including the 3 most significant components. This is seen by the concave line plot for the test data above.

- Use the test set to decide which of the 5 models above is best to predict out of sample. How does this model compare to the previous models you've fit? What are the interpretations of this model's coefficients?

From the above, we find that it is best to select the model with 3 principle components. This model is much worse than our previous models, with a R^2 for the test data set of .177. There are no obvious interpretations of this model's coefficients, other than the weightings of each principle component vector. These weightings are [-34.95509138 -32.72310589 48.22779733] respectively. Further, we have a very large positive intercept of 4598.447129909366.

3.7 Part (i): Beyond Squared Error

We have seen in class that the multiple linear regression method optimizes the Mean Squared Error (MSE) on the training set. Consider the following alternate evaluation metric, referred to as the Root Mean Squared Logarithmic Error (RMSLE):

$$\sqrt{\frac{1}{n} \sum_{i=1}^n (\log(y_i + 1) - \log(\hat{y}_i + 1))^2}.$$

The *lower* the RMSLE the *better* is the performance of a model. The RMSLE penalizes errors on smaller responses more heavily than errors on larger responses. For example, the RMSLE penalizes a prediction of $\hat{y} = 15$ for a true response of $y = 10$ more heavily than a prediction of $\hat{y} = 105$ for a true response of 100, though the difference in predicted and true responses are the same in both cases.

This is a natural evaluation metric for bike share demand prediction, as in this application, it is more important that the prediction model is accurate on days where the demand is low (so that the few customers who arrive are served satisfactorily), compared to days on which the demand is high (when it is less damaging to lose out on some customers).

The following code computes the RMSLE for you:

```
In [36]: #----- rmsle
# A function for evaluating Root Mean Squared Logarithmic Error (RMSLE)
# of the linear regression model on a data set
# Input:
#     y_test (n x 1 array of response variable vals in testing data)
#     y_pred (n x 1 array of response variable vals in testing data)
# Return:
#     RMSLE (float)

def rmsle(y, y_pred):
    # Evaluate squared error, against target labels
    # rmsle = \sqrt{1/n \sum_i (\log(y[i]+1) - \log(y_pred[i]+1))^2}
    rmsle_ = np.sqrt(np.mean(np.square(np.log(y+1) - np.log(y_pred+1))))

    return rmsle_
```

Use the above code to compute the training and test RMSLE for the polynomial regression model you fit in Part (g).

You are required to develop a strategy to fit a regression model by optimizing the RMSLE on the training set. Give a justification for your proposed approach. Does the model fitted using your approach yield lower train RMSLE than the model in Part (g)? How about the test RMSLE of the new model?

Note: We do not require you to implement a new regression solver for RMSLE. Instead, we ask you to think about ways to use existing built-in functions to fit a model that performs well on RMSLE. Your regression model may use the same polynomial terms used in Part (g).

```
In [37]: # your code here
yTest = testBinaryDf["count"]

yHatTest = multModel.predict(multTestPredict)
print("RMSLE for the model in part g: {}".format(rmsle(yTest, yHatTest)))
```

RMSLE for the model in part g: 0.5242723327381363

```
/Applications/anaconda/lib/python3.6/site-packages/ipykernel_launcher.py:13: RuntimeWarning: invalid  
del sys.path[0]
```

```
In [38]: X = multTrainPredict.copy()
        Y = yTrain

        newModel = sm.OLS(list(Y), X).fit()
        yHat = newModel.predict(X)

        error = np.abs(yHat - Y)

        Y = Y.reindex([error.argsort()])
        X = X.reindex([error.argsort()])

        X_aug = X
        Y_aug = Y

        for _ in range(400):
            X_aug = X_aug.append(X[:50])
            Y_aug = Y_aug.append(Y[:50])

        newModel = sm.OLS(list(Y_aug), X_aug).fit()
        yHat = newModel.predict(multTestPredict)
        rmsleVal = rmsle(yTest, yHat)
        print("Our new RMSLE is: {}".format(rmsleVal))
```

```
Our new RMSLE is: 0.5051747912721489.
```

```
/Applications/anaconda/lib/python3.6/site-packages/ipykernel_launcher.py:13: RuntimeWarning: invalid  
del sys.path[0]
```

3.8 Answers to Above Questions

- You are required to develop a strategy to fit a regression model by optimizing the RMSLE on the training set. Give a justification for your proposed approach. Does the model fitted using your approach yield lower train RMSLE than the model in Part (g)? How about the test RMSLE of the new model?

Our strategy to fit a regression model to optimize RMSLE on the training set was the following. We recognized that data points with small absolute errors in an initial fitting of our training data set would decrease the RMSLE. Therefore, we augmented our training data by including data points with small absolute residuals multiple times. This approach intuitively makes sense, as small $|\hat{Y} - Y| \Rightarrow \ln(\hat{Y} + 1) - \ln(Y + 1)$ will also be small, decreasing our overall RMSLE.

Our new model has an *RMSLE* value on the test data of .505, which is better than our previous model from part (g)'s score of .524.

3.9 Part (j): Dealing with Erroneous Labels

Due to occasional system crashes, some of the bike counts reported in the data set have been recorded manually. These counts are not very unreliable and are prone to errors. It is known that roughly 5% of the labels in the training set are erroneous (i.e. can be arbitrarily different from the true counts), while all the labels in the test set were confirmed to be accurate. Unfortunately, the identities of the erroneous records in the training set are not available. Can this information about presence of 5% errors in the training set labels (without details about the specific identities of the erroneous rows) be used to improve the performance of the model in Part (g)? Note that we are interested in improving the R^2 performance of the model on the test set (not the training R^2 score).

As a final task, we require you to come up with a strategy to fit a regression model, taking into account the errors in the training set labels. Explain the intuition behind your approach (we do not expect a detailed mathematical justification). Use your approach to fit a regression model on the training set, and compare its test R^2 with the model in Part (g).

Note: Again, we do not require you to implement a new regression solver for handling erroneous labels. It is sufficient that you to come up with an approach that uses existing built-in functions. Your regression model may use the same polynomial terms used in Part (g).

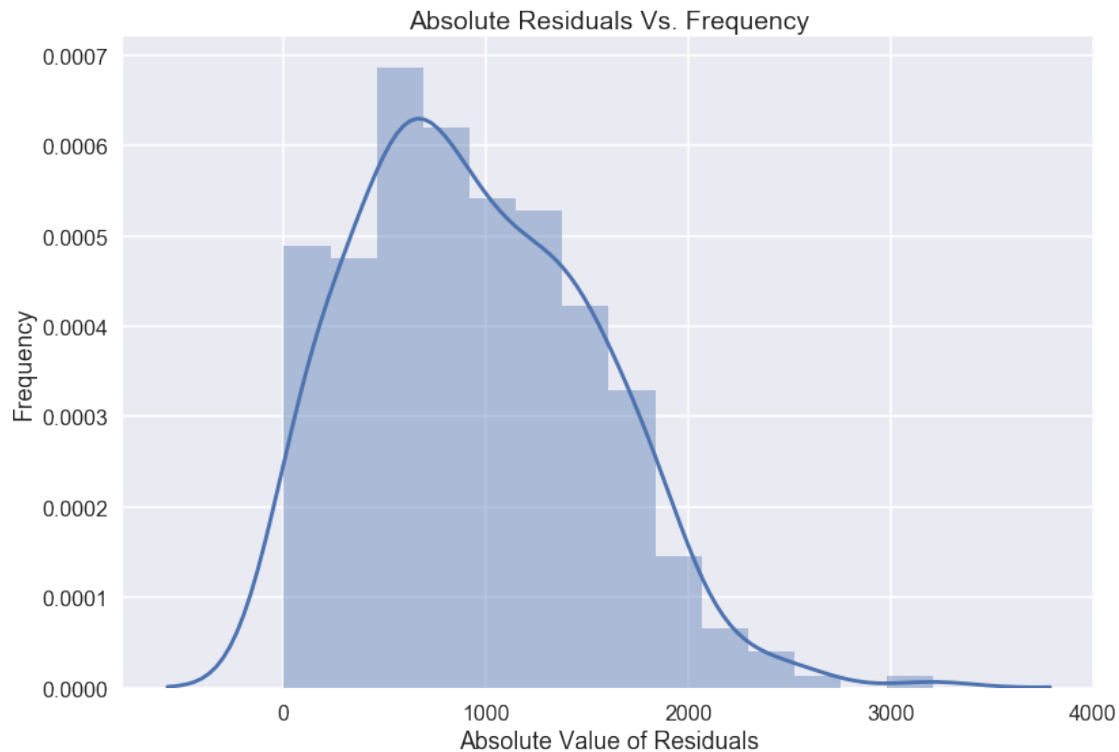
In [39]: *# identify absolute error*

```
multModel = sm.OLS(yTrain, multTrainPredict).fit()
yHatTrain = multModel.predict(multTrainPredict)

errors = yTrain - yHatTrain
abserrors = np.array(abs(errors))

fig, ax = plt.subplots(1)
ax.set_title("Absolute Residuals Vs. Frequency")
ax.set_ylabel("Frequency")
ax.set_xlabel("Absolute Value of Residuals")
sns.distplot(abserrors)
```

Out [39]: <matplotlib.axes._subplots.AxesSubplot at 0x111dc1978>



```
In [40]: dropIndices = abserrors.argsort()[-int(len(abserrors)*0.05):][::-1]

adjYTrain = yTrain.drop(yTrain.index[dropIndices])
adjXTrain = multTrainPredict.drop(multTrainPredict.index[dropIndices])
adjMultModel = sm.OLS(adjYTrain, adjXTrain).fit()
adjMultModel.summary()
```

```
Out[40]: <class 'statsmodels.iolib.summary.Summary'>
"""
                                OLS Regression Results
=====
Dep. Variable:                  count    R-squared:                0.718
Model:                            OLS    Adj. R-squared:           0.678
Method:                 Least Squares    F-statistic:                17.93
Date:                Tue, 10 Oct 2017    Prob (F-statistic):        3.05e-55
Time:                  13:29:30    Log-Likelihood:            -2625.6
No. Observations:                315    AIC:                       5331.
Df Residuals:                    275    BIC:                       5481.
Df Model:                        39
Covariance Type:                nonrobust
=====
                                coef    std err          t      P>|t|      [0.025    0.975]
-----
-----
```

holiday	-89.0664	363.130	-0.245	0.806	-803.935	625.802
workingday	85.7392	152.585	0.562	0.575	-214.643	386.122
temp_norm	753.7641	698.454	1.079	0.281	-621.233	2128.761
atemp_norm	931.9529	655.442	1.422	0.156	-358.368	2222.274
humidity_norm	-695.2858	144.789	-4.802	0.000	-980.322	-410.250
windspeed_norm	-443.0108	138.341	-3.202	0.002	-715.353	-170.669
weather_1.0	1007.8855	509.110	1.980	0.049	5.638	2010.133
weather_2.0	1082.5915	463.259	2.337	0.020	170.607	1994.576
day_of_week_0.0	-340.7946	232.970	-1.463	0.145	-799.425	117.836
day_of_week_1.0	-202.8532	149.029	-1.361	0.175	-496.235	90.529
day_of_week_2.0	-207.6025	172.805	-1.201	0.231	-547.791	132.586
day_of_week_3.0	138.3122	184.038	0.752	0.453	-223.990	500.615
day_of_week_4.0	112.0110	185.816	0.603	0.547	-253.791	477.813
day_of_week_5.0	156.8052	172.188	0.911	0.363	-182.168	495.779
season_1.0	-1662.8420	429.543	-3.871	0.000	-2508.452	-817.232
season_2.0	-675.8582	493.064	-1.371	0.172	-1646.518	294.802
season_3.0	226.2033	385.309	0.587	0.558	-532.326	984.733
month_2.0	-248.0666	388.167	-0.639	0.523	-1012.224	516.091
month_3.0	-440.3675	425.067	-1.036	0.301	-1277.166	396.430
month_4.0	-492.6351	603.950	-0.816	0.415	-1681.588	696.318
month_5.0	-1301.7164	635.376	-2.049	0.041	-2552.536	-50.897
month_6.0	-1853.3696	656.642	-2.822	0.005	-3146.053	-560.686
month_7.0	-1693.9734	701.678	-2.414	0.016	-3075.316	-312.631
month_8.0	-2041.3922	699.262	-2.919	0.004	-3417.979	-664.805
month_9.0	-1316.2651	617.398	-2.132	0.034	-2531.692	-100.838
month_10.0	-1272.8666	574.251	-2.217	0.027	-2403.353	-142.380
month_11.0	-960.2166	550.751	-1.743	0.082	-2044.441	124.008
month_12.0	-618.9551	447.518	-1.383	0.168	-1499.952	262.042
const	6165.5547	779.588	7.909	0.000	4630.837	7700.273
temp_norm_2	-2178.4066	754.456	-2.887	0.004	-3663.650	-693.163
atemp_norm_2	1441.5058	728.145	1.980	0.049	8.060	2874.952
humidity_norm_2	-72.8165	143.386	-0.508	0.612	-355.090	209.457
windspeed_norm_2	-70.4969	116.980	-0.603	0.547	-300.787	159.793
temp_norm_3	28.2302	252.891	0.112	0.911	-469.618	526.078
atemp_norm_3	-327.9724	225.474	-1.455	0.147	-771.847	115.903
humidity_norm_3	-22.5567	41.820	-0.539	0.590	-104.886	59.772
windspeed_norm_3	54.3955	60.873	0.894	0.372	-65.442	174.233
temp_norm_4	-2.6787	157.059	-0.017	0.986	-311.868	306.511
atemp_norm_4	-51.0898	135.467	-0.377	0.706	-317.775	215.595
humidity_norm_4	-23.4383	28.842	-0.813	0.417	-80.217	33.341
windspeed_norm_4	-18.2169	27.861	-0.654	0.514	-73.065	36.631
=====						
Omnibus:	53.297	Durbin-Watson:	1.956			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	12.734			
Skew:	-0.055	Prob(JB):	0.00172			
Kurtosis:	2.021	Cond. No.	3.14e+16			
=====						

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified
[2] The smallest eigenvalue is 3.81e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.
"""
```

```
In [41]: yHatAdjTest = adjMultModel.predict(multTestPredict)
        r2_score(yTest, yHatAdjTest)
```

```
Out[41]: 0.26553825406157139
```

3.10 Answer to Above Questions

- As a final task, we require you to come up with a strategy to fit a regression model, taking into account the errors in the training set labels. Explain the intuition behind your approach (we do not expect a detailed mathematical justification). Use your approach to fit a regression model on the training set, and compare its test R^2 with the model in Part (g).

Our strategy for the above problem was the following. We wanted to find the values in our training data set where the absolute value of their residuals were particularly large. By excluding the top 5% largest absolute residuals, we ensure that any erroneous entries are not outliers skewing our model unfavorably. We recognize that there is a chance we have not removed exactly all of the errors, but in the event that the errors are clustered with the majority of our correct data, our model should still accurately predict our training data set.

This hypothesis is supported by the fact that our R^2 on the test data set did not significantly decrease from our previous model with the outliers in the training data set. Our new R^2 is .266, which is slightly lower than .277 from our model in part g.

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```